

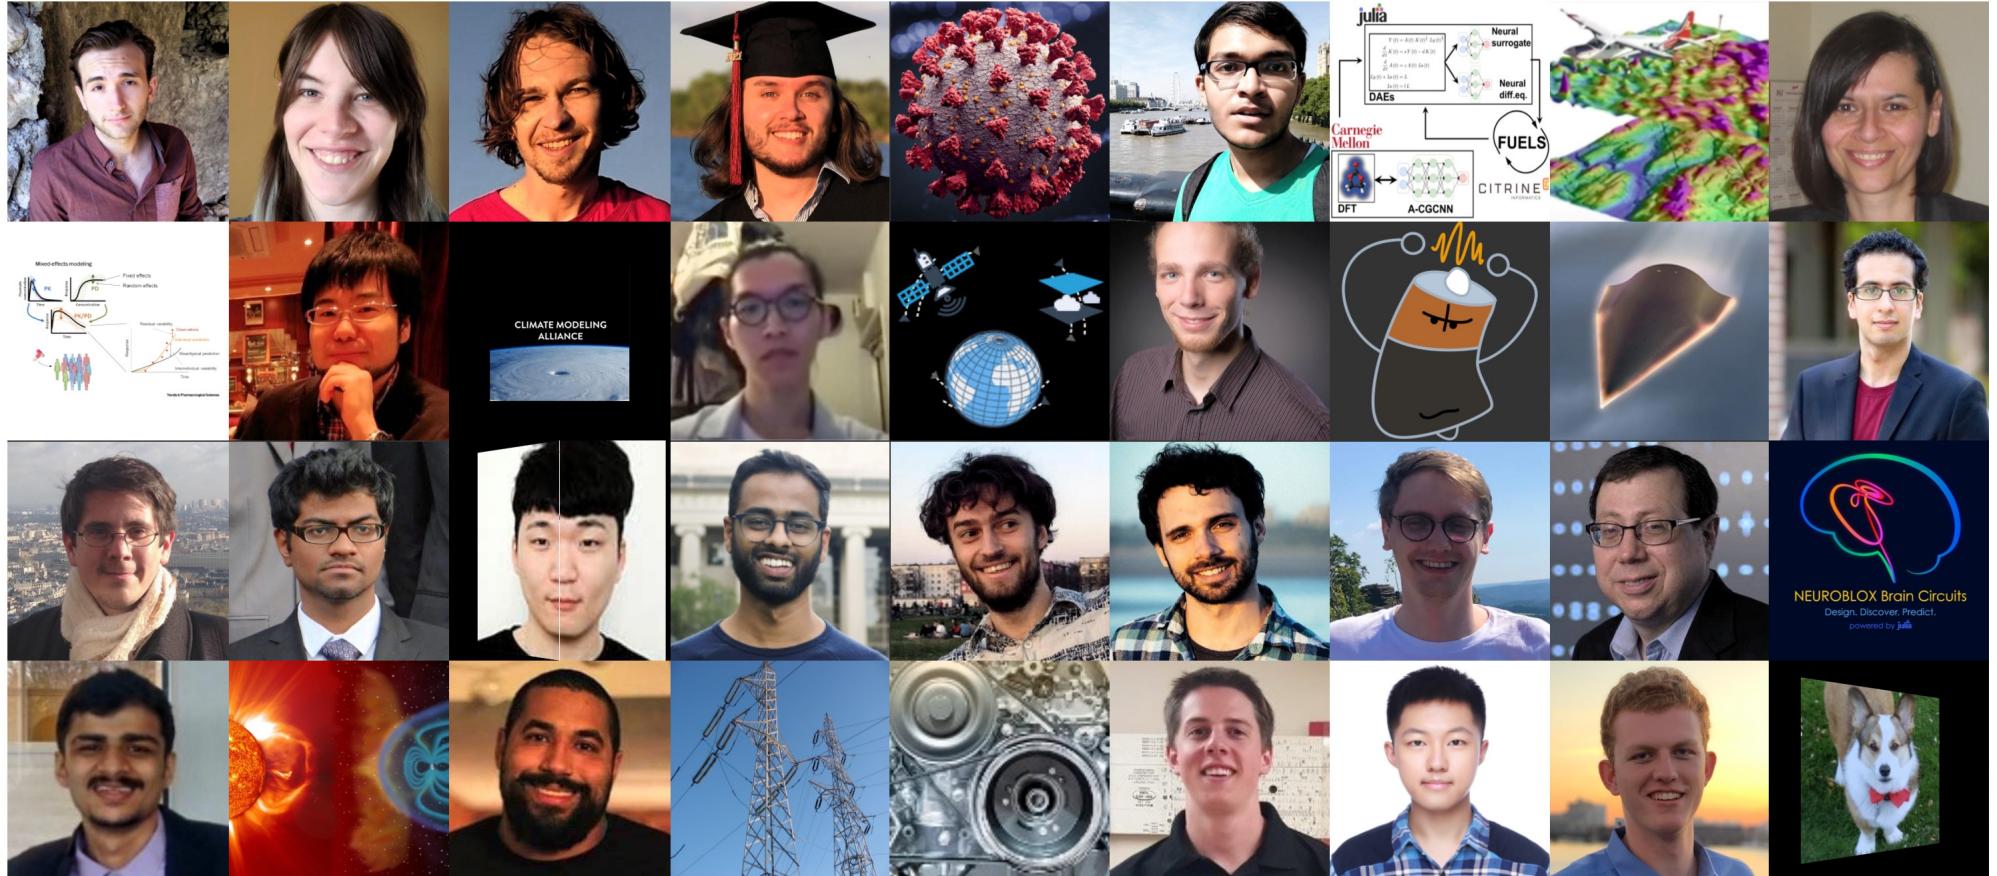
Parallel Computing's Biggest Challenge



Alan Edelman
SIAM PP 2022
February 23, 2022



MIT Julialab: julia.mit.edu



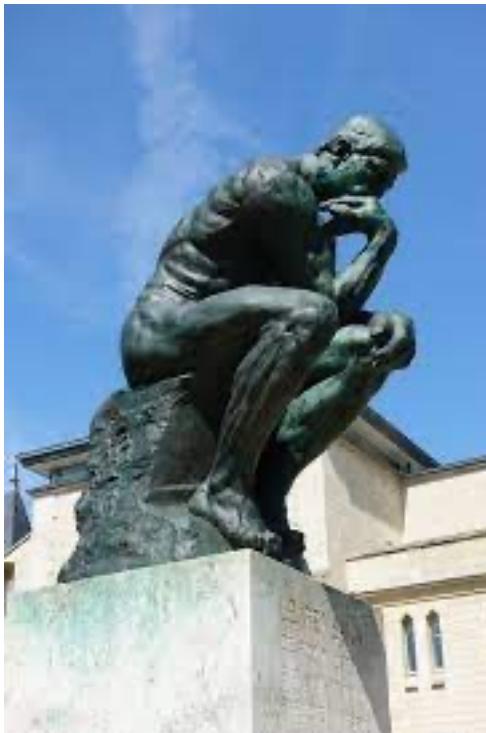
TEDx A programming language to heal the planet together...

Modeling Spacecraft Separation Dynamics in Julia...

THE JULIA SCIML ECOSYSTEM: SCIENTIFIC MACHINE LEARNING AS A SOFTWARE PROBLEM

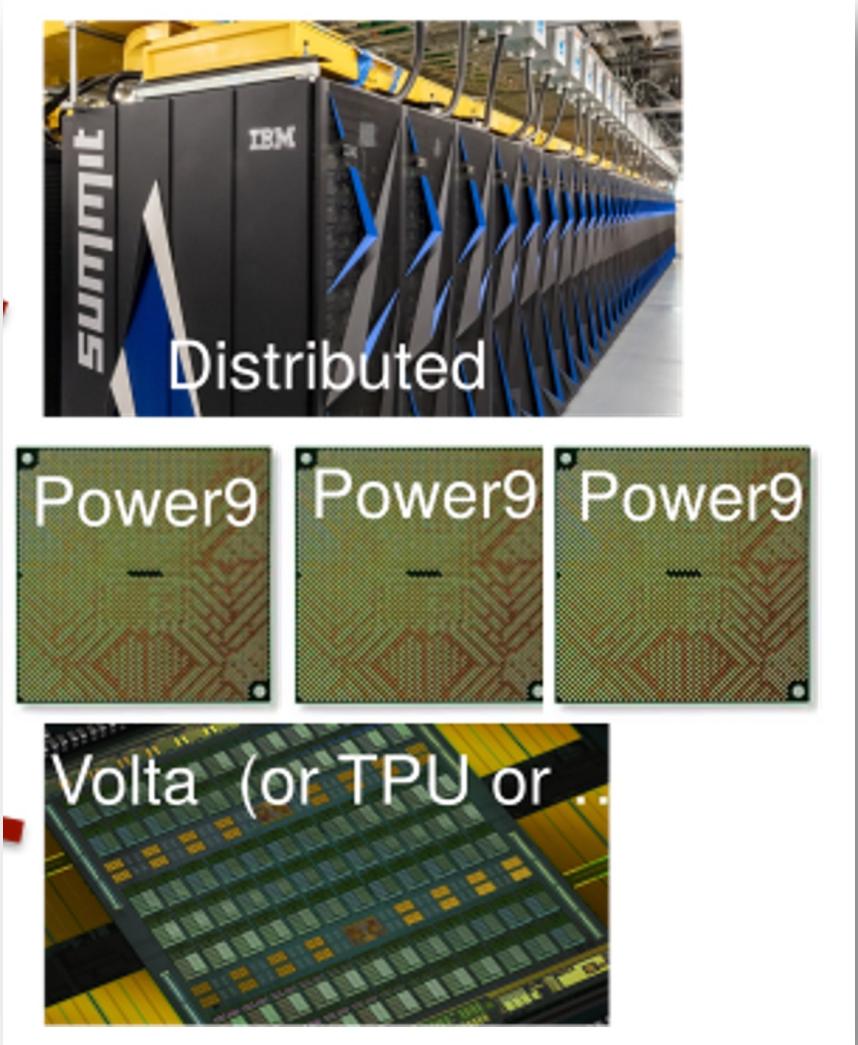
Audience Participation

SIAM Conference on Parallel Processing for Scientific Computing (PP22)



Biggest
Challenge in
our field?

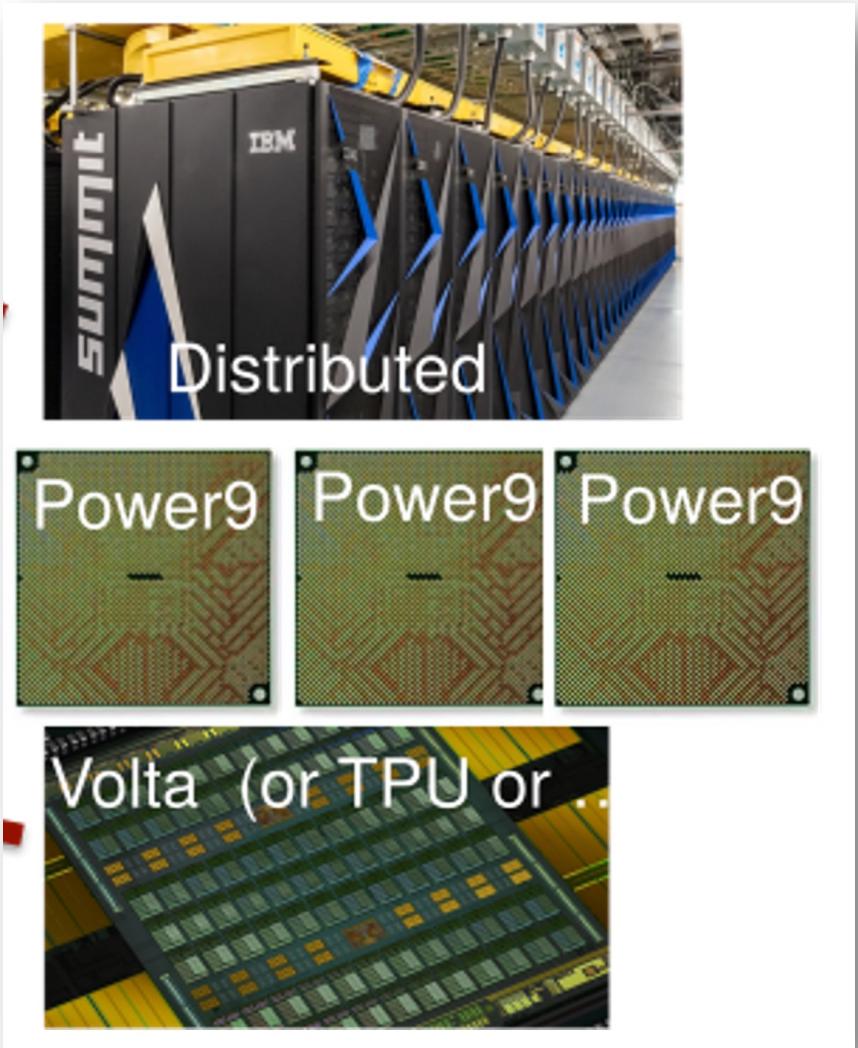
Heterogeneity Some may say ...



but what is usually
meant by this?

Of course Performance!!!

Heterogeneity Some may say ...



but should it be?

Heterogeneity Some may say ...

How can we program these complex systems for performance?

Heterogeneity Some may say ...

WARNING

How can we program these complex systems for **performance?**

Optimizing solely for
performance leads to decades of
missed opportunities

NOTHING
IS MORE
EXPENSIVE
THAN A MISSED
OPPORTUNITY

H. Jackson Brown, Jr

Heterogeneity Some may say ...

WARNING

How can we program these complex systems for **performance**?

If I could influence the field, I would **remove performance** as the sole metric of success.

NOTHING
IS MORE
EXPENSIVE
THAN A MISSED
OPPORTUNITY

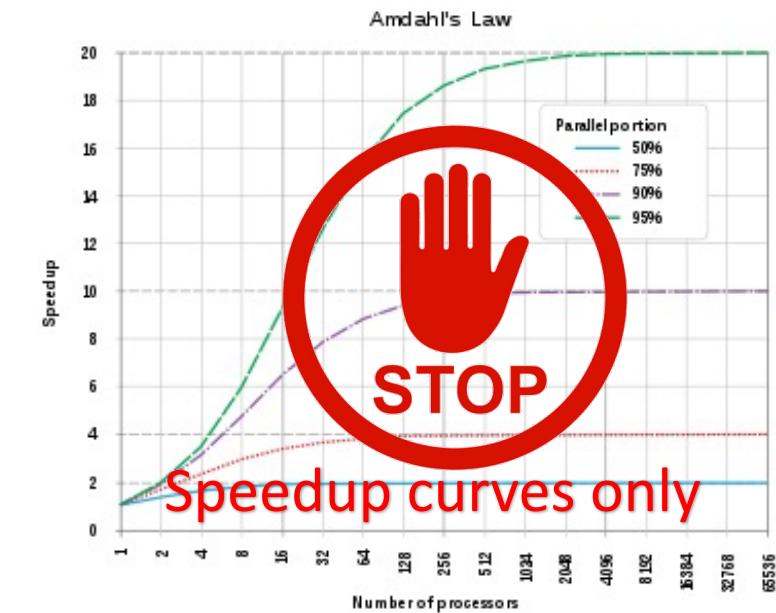
H. Jackson Brown, Jr

Heterogeneity Some may say ...

WARNING

How can we program these complex systems for performance?

If I could influence the field, I would **remove performance** as the sole metric of success.



Heterogeneity Some may say ...

WARNING

How can we program these complex systems for **performance**?

I would replace (with discretion)
with the size of the user or
developer community. (Impact!)

Performance is but one component!

Performance indicator

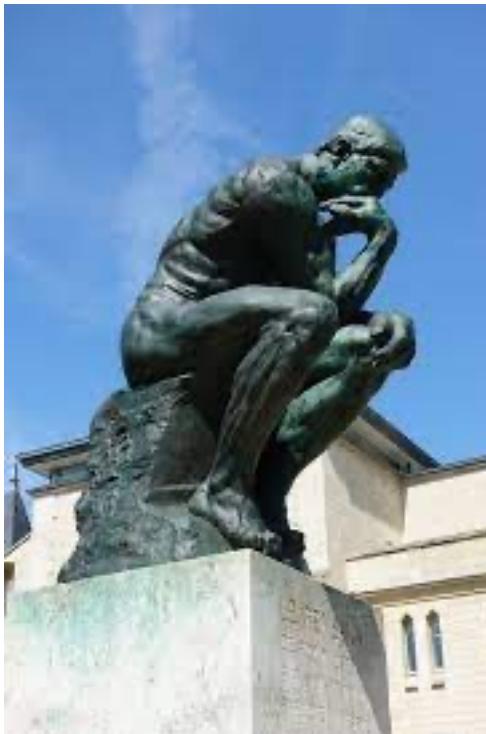
From Wikipedia, the free encyclopedia

Contents [hide]

- 1 Categorization of indicators
- 2 Points of measurement
- 3 Identifying indicators of organization
- 4 Examples
 - 4.1 Accounts
 - 4.2 Marketing and sales
 - 4.3 Manufacturing
 - 4.4 Professional Services
 - 4.5 System operations
 - 4.6 Project execution
 - 4.7 Supply chain management
 - 4.8 Government
 - 4.9 Further performance indicators

Audience Participation

SIAM Conference on Parallel Processing for Scientific Computing (PP22)



Biggest
Challenge in
our field?

Time for a summary

- Performance alone has lead to opportunity loss
- performance / productivity / reusability / community /...
- What technically is desirable in the bedrock to achieve the impact we deserve?
- More and more are getting the memo:
- Especially the younger crowd



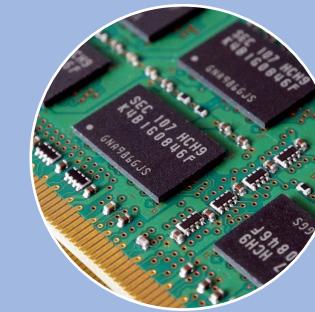
What about programming models?



Data parallel – often associated with **GPU**

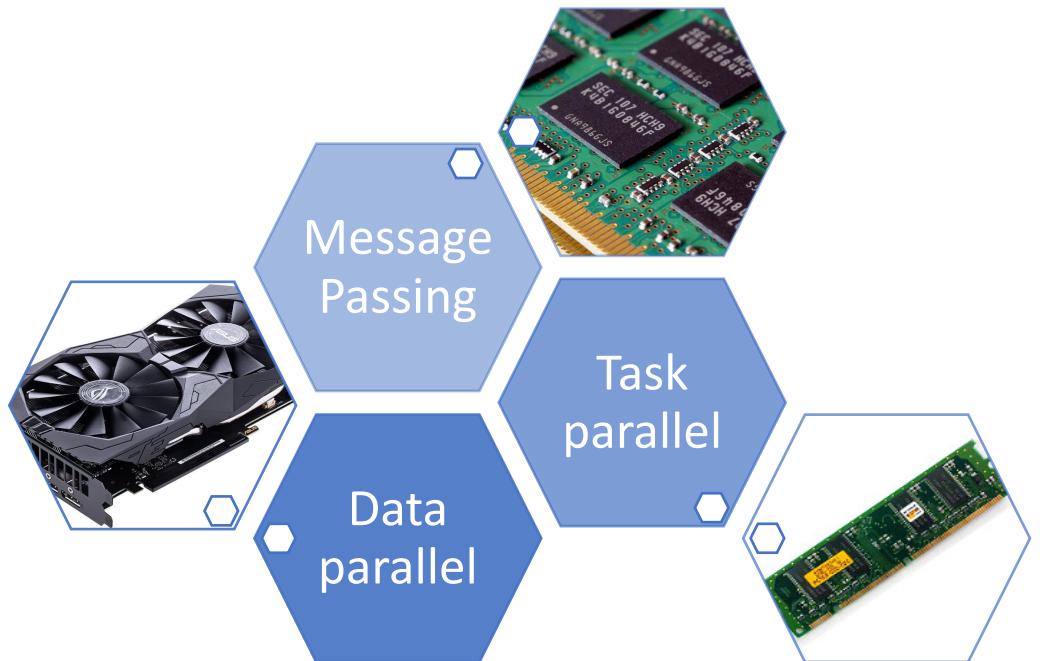


Task parallel – often associated with
Shared Mem



Message Passing – often associated with
Distributed Mem

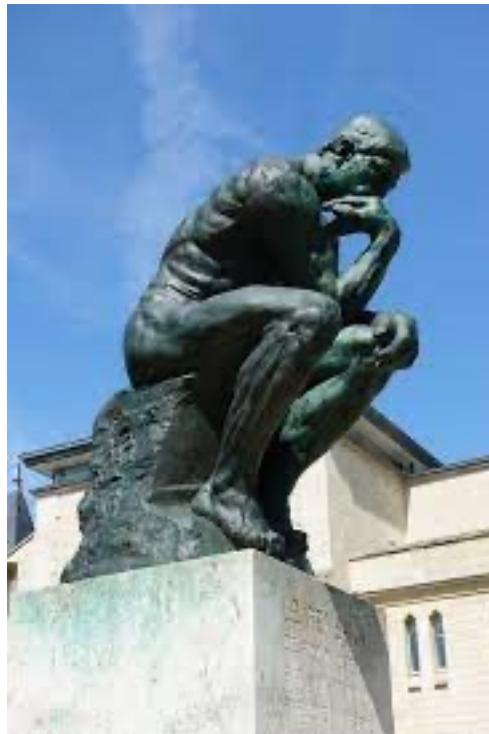
What about programming models?



- Each has their own style
- Difficult to composite due to different styles
- Lots of inconsistencies
- Fragmented communities?
- Difficult to reuse and extend existing work

Audience Participation

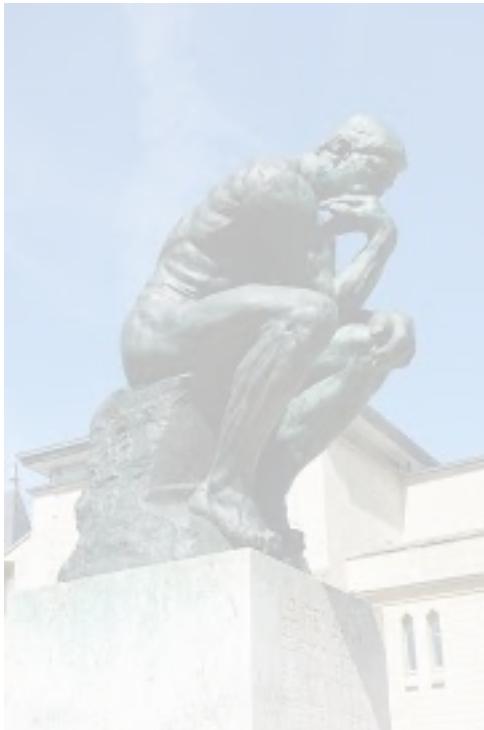
SIAM Conference on Parallel Processing for Scientific Computing (PP22)



Biggest
Challenge in
our field?

Audience Participation

SIAM Conference on Parallel Processing for Scientific Computing (PP22)



• **Language!** Biggest Challenge in our field?



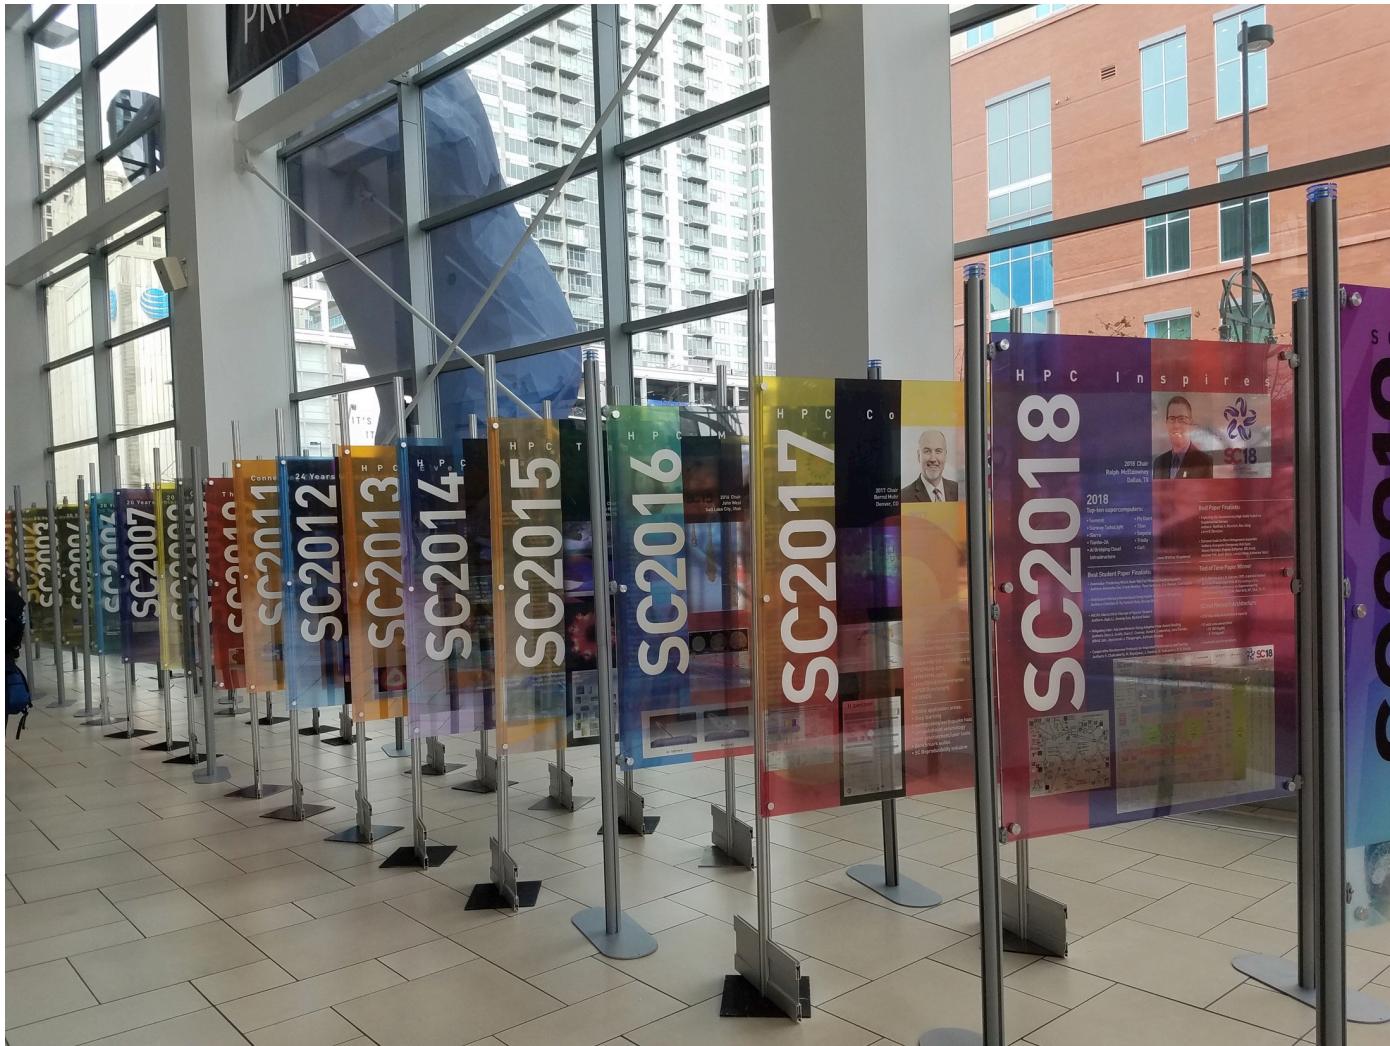
1988

Programming Languages (x2005)

2010

2019

The dark decade



Something seems to have gone wrong!

What did we miss 20 years ago?

Chapel, Fortress and X10:
novel languages for HPC

1 Introduction

Michèle Weiland

In 2002, DARPA¹ launched their High Productivity Computing Systems (HPCS) programme, offering funding for industry and academia to research the development of computing systems which focus on high productivity along with high performance. Part of this programme concentrates on the specification of novel languages for the HPC community. In Phase 2 of the programme (July 2003 - July 2006) the three remaining partners that were awarded funding were Cray, IBM and SUN²; SUN were eventually dropped from the programme at the start of Phase 3.[1]

This report will introduce the novel programming languages *Chapel* (Cray), *Fortress* (SUN) and *X10* (IBM) that were developed as part of the HPCS programme.

10 years ago

Why We Created Julia

14 February 2012 | Jeff Bezanson Stefan Karpinski Viral B. Shah Alan Edelman



[Jeff Bezanson Stefan Karpinski Viral B. Shah Alan Edelman](#)

In short, because we are greedy.

We are power Matlab users. Some of us are Lisp hackers. Some are Pythonistas, others Rubyists, still others Perl hackers. There are those of us who used Mathematica before we could grow facial hair. There are those who still can't grow facial hair. We've generated more R plots than any sane person should. C is our desert island programming language.

We love all of these languages; they are wonderful and powerful. For the work we do — scientific computing, machine learning, data mining, large-scale linear algebra, distributed and parallel computing — each one is perfect for some aspects of the work and terrible for others. Each one is a trade-off.

We are greedy: we want more.

We want a language that's open source, with a liberal license. We want the speed of C with the dynamism of Ruby. We want a language that's homoiconic, with true macros like Lisp, but with obvious, familiar mathematical notation like Matlab. We want something as usable for general programming as Python, as easy for statistics as R, as natural for string processing as Perl, as powerful for linear algebra as Matlab, as good at gluing programs together as the shell. Something that is dirt simple to learn, yet keeps the most serious hackers happy. We want it interactive and we want it compiled.

(Did we mention it should be as fast as C?)

To solve the parallel computing problem – and it wouldn't be easy.

We'd need help from **the community**.

Today

Keno Fischer (@keno)

Ten years ago, when "Why we created Julia" was published, I was just finishing an exchange student at a tiny school on the Eastern Shore of Maryland. With school attempting school projects and online courses in Octave, college acceptance in mind, resonated strongly with the need for better tools in computational science and I worked on everything from astronomy, over homomorphic encryption, to simulation in between. There have been many memorable events along the way. For example, I am grateful and humbled. Julia is taking big strides now, since its first appearance 10 years ago. I am really excited about its future and honored to be a part of it.

Avik Sengupta (@aviks)

My first reaction on seeing the headline on Hacker News that day was a decidedly lukewarm "oh, do we really need yet another programming language". Reading the blog post however piqued my interest – "it surely can't deliver on all that it promises, can it?". It was a slow day at work, so I downloaded the source, and was surprised to see that it built successfully on the first try. My [first PR](#) came two weeks later. In adding a new

computing was certainly possible. In the , as well as being able to write Julia

Jerry Ling (@Moelf)

I was, like many people, moderately skeptical when I saw Julia's promise of "the best of both worlds" the first time; in fact, I remember walking away not particularly impressed the first time I tried Julia in 2018, largely due to latency (ITTF) issue with plotting. That soon changed when I tried again near v1.0 release after hearing a PhD student (Katharine Hyatt) in UCSB physics was doing cool quantum stuff in Julia. Things moved fast after that: once you know the two-language problem, it's everywhere in physics. So many things are "free" in Julia, Python can be skipped because Julia is not slow, auto diff would have helped with problem in physics etc. Today, though not as proficient in compiler or typing system than I was back then, I strive to bring better computing tools to physics, and I hope to witness a rapid next decade.

Cristóvão Sousa (@cdsousa)

Julia is the language I was longing for. And in the late summer of 2013, that "Why We Created Julia" was mind-blowing to me! I had been taught to program in C and C++ and later I learned some MATLAB too. I always looked for performance and I'm the kind of person that prefers to waste more time optimizing code than wasting little time waiting for it to run. But on the other hand, I need to write math/scientific code and I like to do it in a high-level language. When starting the PhD in 2009, along with C++ I wanted to use an open-source higher-level language, so I opted for Python. However, the slowness was a bummer and I had to rely on workarounds like Cython. I was not happy, and I was greedy, always wondering "Why can't a better language, high-level but performant, exist?". In the summer of 2013, I found Julia. It was exactly "Oh no! Yet another language?!" but it was "WOW! YAY!". It's been 10 years since then and I'm still very happy with Julia.

Chris Rackauckas (@chrisrackauckas)

I came rather late to the scene, around 2016. I had written GUIs in R, libraries in Python and MATLAB, handled MEX files and wrote my MPI using C. I wrote some ODE solvers using Fortran like a good ol' kid. I did my due diligence. It was an absolute mess, especially being a Windows user, and so when I found Julia I was astonished to find something that would actually work. No more SciPy installation error from some compiler

missing on Windows: Julia was the compiler. And more and more, Julia is the compiler. I

Julian Samaroo (@jpsamaroo)

About 8 years ago, I was a student studying neuroscience in college, very interested in computational modeling. I had implemented some basic spiking neural network simulators in MATLAB, and ported them to JavaScript, but always found both languages lacking. Still, I persisted with JavaScript, until about 3 years later, when I first heard about Julia from somewhere on the interwebs. Out of boredom and a mild displeasure with JavaScript, I decided to give Julia a try. I will be honest; I had no shortage of problems using it! The errors were unintuitive, the syntax was strange, and it was *slow* (to compile, but what was the difference?). I kept using JavaScript for work, but slowly on the side I started getting more familiar with Julia, and tried using it for some of my simulations and side projects. Skipping to about 4 years ago, I got "nerd-sniped" by a certain Valentin Churavy to implement support for AMD GPUs in Julia, since I was so vocal about wanting to support a more open-source friendly GPU vendor. Over the next 2 years, I developed AMDGPUnative.jl (now AMDGPU.jl) and Dagger.jl in my free time, slowly gaining an appreciation for Julia, and attended my very first JuliaCon in 2019! Somewhat shortly afterwards, Valentin reached out from the JuliaLab and offered me an RSE position (which was an offer I could have only dreamed of at that point in life!), which I gladly accepted. Since then, I've been working full-time on AMDGPU.jl, Dagger.jl, BPFnative.jl, and Julia itself, and have been loving every moment of it. I don't think I could ever give up the wonderful ecosystem, language, and people that I've come to know and love over the last 5 years, and I look forward to many more wonderful years to come.

Matt Brzezinski (@mattbrzezinski)

July 2019 I left my job at AWS and joined Julia. I first heard about the new programming language from a colleague at AWS who had just started using it. I was immediately hooked. Julia is a production environment. I spent my first week at Invenia porting various internal and external packages from Julia 0.7 to 1.0 and found it to be a simple and intuitive language to use coming from a Python background. My second week was spent in Baltimore meeting and chatting and learning with the community in Baltimore for JuliaCon 2019. I was hooked. Almost 3 years later I've found that Julia has been my go-to language for nearly everything. The ease of use, speed and community keep bringing me back for more. I'm excited to see where it goes from here.

Jose Storopoli (@storopoli)

I stumbled into Julia while trying to do a crazy data transformation in R that was taking forever. I was instantaneously hooked! The syntax is so easy, and it has also a "math" feel to it. You can use all your ϵ and δ , along with the statistician's favorites α and β . Since then, I've become a contributor to the [JuliaStats](#) and [Turing](#) (Bayesian modeling) Julia's ecosystems. The community is great and very welcoming. I've made amazing friendships here, co-authored a [free open access and open source Julia Data Science book](#) with Rik Huijzer and Lazaro Alonso. The book has been translated to Portuguese and Chinese by volunteers, that which I am grateful and humbled. Julia is taking big strides now, since its first appearance 10 years ago. I am really excited about its future and honored to be a part of it.

To solve the parallel computing problem – and it wouldn't be easy.

We'd need help from **the community**.

Why We Use Julia, 10 Years Later

14 February 2022 | The Julia Community



What did we miss 20 years ago?

A multidisciplinary user base for high productivity & high performance

What did we miss 20 years ago?

A multidisciplinary user base for high productivity & high performance

multiple dispatch
static + dynamic
type specialization



...

What did we miss 20 years ago?

A multidisciplinary user base for high productivity & high performance

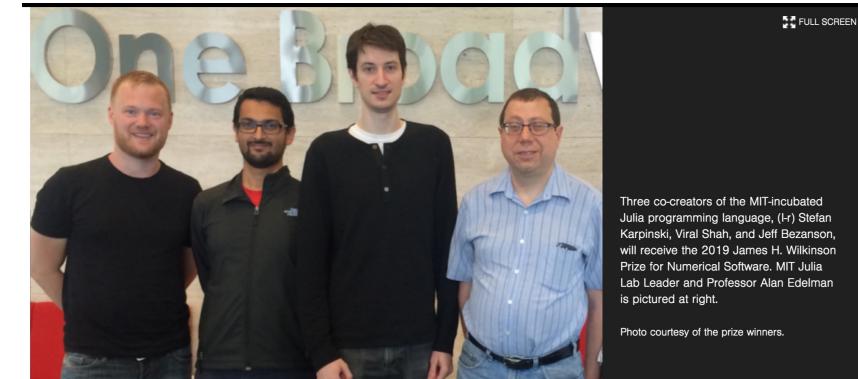
multiple dispatch
static + dynamic
type specialization

...



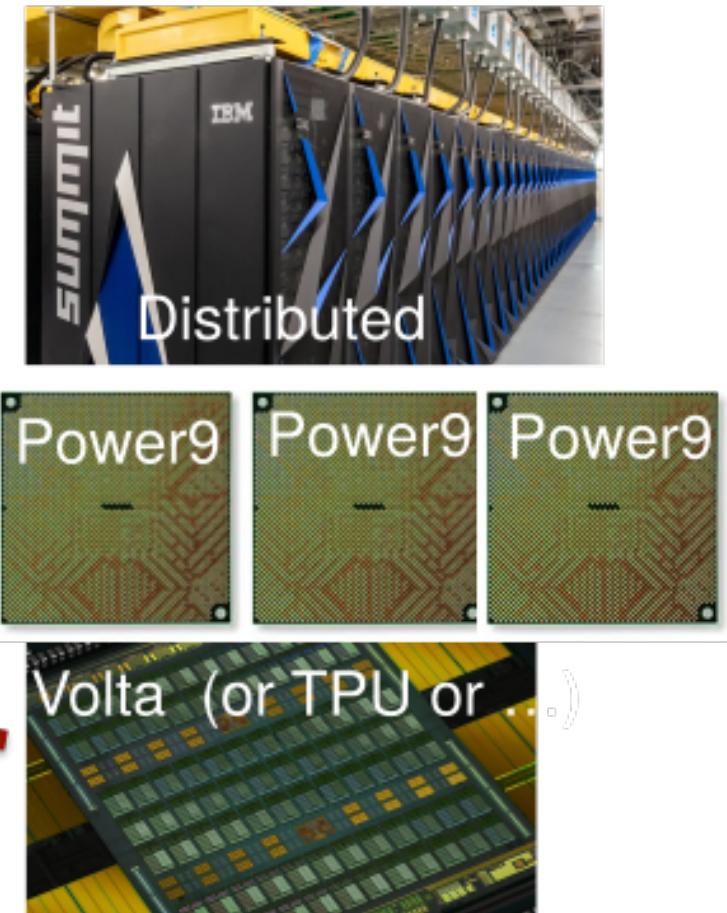
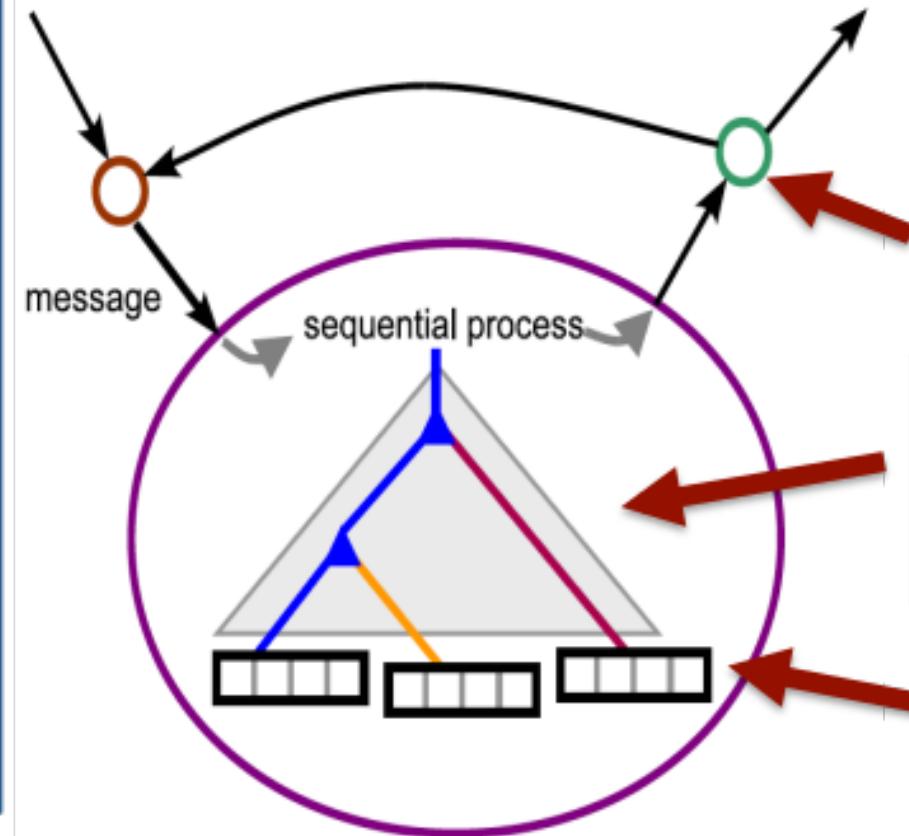
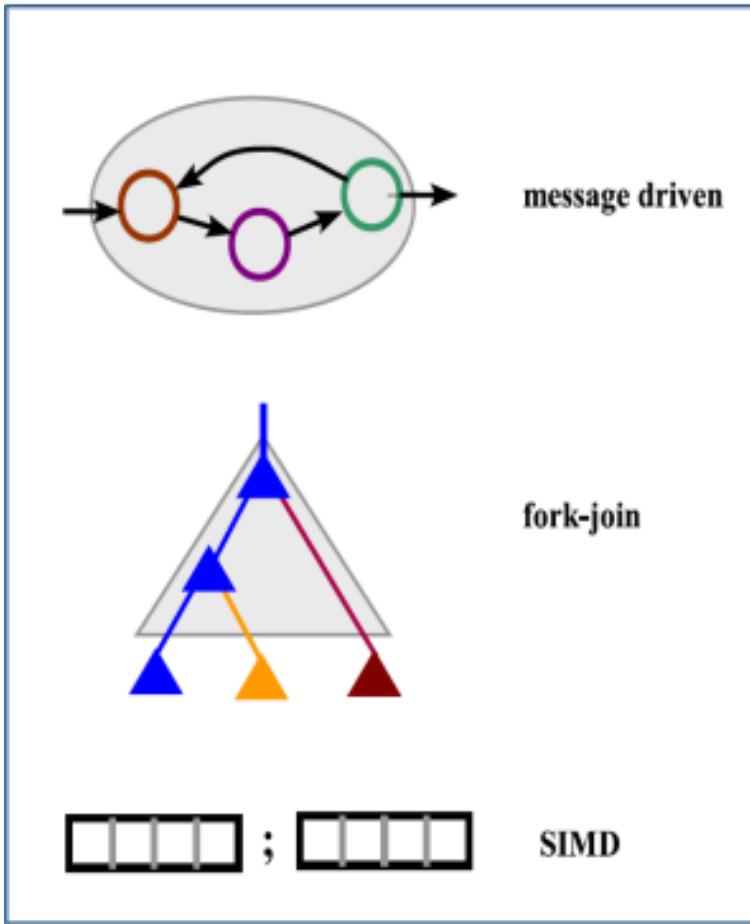
“Patience is a virtue.”

We knew the right user base starts serial and takes over 10 years to mature.



Julia language co-creators win James H. Wilkinson Prize for Numerical Software

Back to Heterogeneity: Towards Higher Level Abstractions



Julia HPC

Accelerators



AMDGPU.jl



CUDA.jl



OneAPI.jl

JuliaGPU/ GPUArrays.jl



Reusable array functionality for Julia's various GPU backends.

<https://github.com/JuliaGPU/KernelAbstractions> ::

KernelAbstractions.jl - Heterogeneous programming in Julia

Heterogeneous programming in Julia. Contribute to JuliaGPU/KernelAbstractions.jl development by creating an ... [JuliaGPU / KernelAbstractions.jl Public](#).

Shared Mem

JuliaFolds/FLoops.jl

Fast sequential, threaded, and distributed for-loops for Julia—fold for humans™



Announcing composable multi-threaded parallelism in Julia

23 July 2019 | Jeff Bezanson (Julia Computing), Jameson Nash (Julia Computing), Kiran Pamnany (Intel)

Base / Multi-Threading

Multi-Threading

[Base.Threads.@threads](#) – Macro

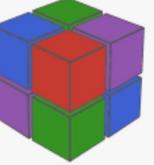
Julia Atomics Manifesto

This proposal aims to define the memory model of Julia and to provide certain guarantees in the presence of data races, both by default and through providing intrinsics to allow the user to specify the level of guarantees required. This should allow native implementation in Julia of simple system primitives (like mutexes), interoperate with native system code, and aim to give generally explainable behaviors without incurring significant performance cost. Additionally, it strives to be general-purpose and yet clear about the user's intent—particularly with respect to ensuring that an atomic-type field is accessed with proper care for synchronization.

Distributed

JuliaParallel/MPI.jl

MPI wrappers for Julia

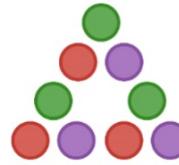


GitHub - eth-cscs/ImplicitGlobalGrid.jl ...
github.com



JuliaParallel/ Dagger.jl

A framework for out-of-core and parallel execution



Standard Library / Distributed Computing

Distributed Computing

Julia HPC

Phase 1: Building up these tools

Phase 2: Letting these tools **break out** from these separate categories

Accelerators



AMDGPU.jl



CUDA.jl



OneAPI.jl

JuliaGPU/ GPUArrays.jl



Reusable array functionality for Julia's various GPU backends.

<https://github.com/JuliaGPU/KernelAbstractions> ::

KernelAbstractions.jl - Heterogeneous programming in Julia

Heterogeneous programming in Julia. Contribute to JuliaGPU/KernelAbstractions.jl development by creating an ... [JuliaGPU / KernelAbstractions.jl Public](#).

Shared Mem

JuliaFolds/FLoops.jl

Fast sequential, threaded, and distributed for-loops for Julia—fold for humans™



Announcing composable multi-threaded parallelism in Julia

23 July 2019 | Jeff Bezanson (Julia Computing), Jameson Nash (Julia Computing), Kiran Pamnany (Intel)

Base / Multi-Threading

Multi-Threading

`Base.Threads.@threads` – Macro

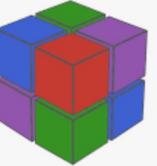
Julia Atomics Manifesto

This proposal aims to define the memory model of Julia and to provide certain guarantees in the presence of data races, both by default and through providing intrinsics to allow the user to specify the level of guarantees required. This should allow native implementation in Julia of simple system primitives (like mutexes), interoperate with native system code, and aim to give generally explainable behaviors without incurring significant performance cost. Additionally, it strives to be general-purpose and yet clear about the user's intent—particularly with respect to ensuring that an atomic-type field is accessed with proper care for synchronization.

Distributed

JuliaParallel/MPI.jl

MPI wrappers for Julia



GitHub - eth-cscs/ImplicitGlobalGrid.jl ...
github.com



JuliaParallel/
Dagger.jl

A framework for out-of-core and parallel execution

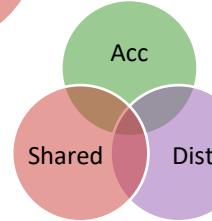
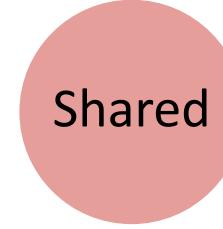


Standard Library / Distributed Computing

Distributed Computing

Julia HPC

Phase 1: Building up these tools



Phase 2: Letting these tools break out from these separate categories



Phase 3: Higher level abstractions

Accelerators



AMDGPU.jl



CUDA.jl



OneAPI.jl

Shared Mem

JuliaFolds/**FLoops.jl**

Fast sequential, threaded, and distributed for-loops
for Julia—fold for humans™



Announcing composable multi-threaded
parallelism in Julia

23 July 2019 | Jeff Bezanson (Julia Computing), Jameson Nash (Julia Computing), Kiran Pamnany (Intel)

Base / Multi-Threading

Multi-Threading

Base.Threads.@threads – Macro

Julia Atomics Manifesto

This proposal aims to define the memory model of Julia and to provide certain guarantees in the presence of data races, both by default and through providing intrinsics to allow the user to specify the level of guarantees required. This should allow native implementation in Julia of simple system primitives (like mutexes), interoperate with native system code, and aim to give generally explainable behaviors without incurring significant performance cost. Additionally, it strives to be general-purpose and yet clear about the user's intent—particularly with respect to ensuring that an atomic-type field is accessed with proper care for synchronization.

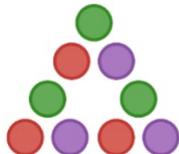
Distributed

JuliaParallel/MPI.jl

MPI wrappers for Julia

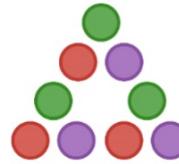


GitHub - eth-cscs/ImplicitGlobalGrid.jl ...
github.com



JuliaParallel/ Dagger.jl

A framework for out-of-core and parallel execution



Standard Library / Distributed Computing

Distributed Computing

JuliaGPU/ GPUArrays.jl



Reusable array functionality for Julia's various GPU
backends.

<https://github.com/JuliaGPU/KernelAbstractions> ::

KernelAbstractions.jl - Heterogeneous programming in Julia

Heterogeneous programming in Julia. Contribute to JuliaGPU/KernelAbstractions.jl development
by creating an ... JuliaGPU / KernelAbstractions.jl Public.

Case study: Histogram $\text{diag}(\text{qr}(\text{randn}(n,n)))$ can't help myself, I ❤ randomized linear algebra

$$H_1 H_2 \cdots H_{n-1} H_n \times \text{randn}(n, n) = R_n \sim$$

H: Householder reflector

G: standard normal

$$\begin{pmatrix} \chi_n & G & G & \dots & G & G & G \\ & \chi_{n-1} & G & \dots & G & G & G \\ & & \chi_{n-2} & \dots & G & G & G \\ & & & \ddots & \vdots & \vdots & \vdots \\ & & & & \chi_3 & G & G \\ & & & & & \chi_2 & G \\ & & & & & & \chi_1 \end{pmatrix}$$

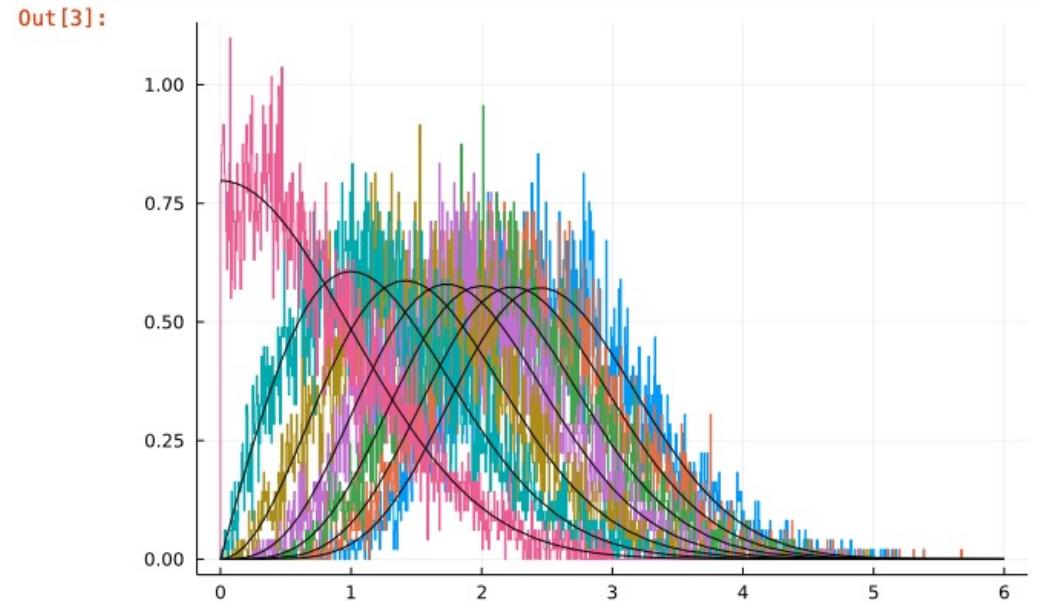
Write Once, Execute Anywhere!

```
function qrhist():
    nsamples = 2^13,
    nbins = 1_000,
    xmin::Real = 0.0,
    xmax::Real = 6.0,
    matsize::Val{m} = Val(7),
    executor = nothing,
) where {m}
    hist = fill!(arrayon(executor) Int, (nbins, m)), 0)
    dx = (xmax - xmin) / nbins
    @floop executor for _i in 1:nsamples
        M = randn(SMatrix{m,m,Float32})
        _q, r = qr(M)
        for i in 1:m
            bin = floor(Int, (abs(r[i, i]) - xmin) / dx) + 1
            @inbounds @atomic :monotonic hist[max(begin, min(bin, end)), i] += 1
        end
    end
    edges = collect(xmin:dx:xmax)
    return [Histogram(edges, collect(h), :left, false) for h in eachcol(hist)]
end
```

Execution on a single CPU

```
In [3]: hist = normalize.(qrhist executor = SequentialEx() nsamples = 2^13))

plt = plot(hist, seriestype = :step, legend = nothing)
plot_chi!(plt, length(hist))
```



Write Once, Execute Anywhere!

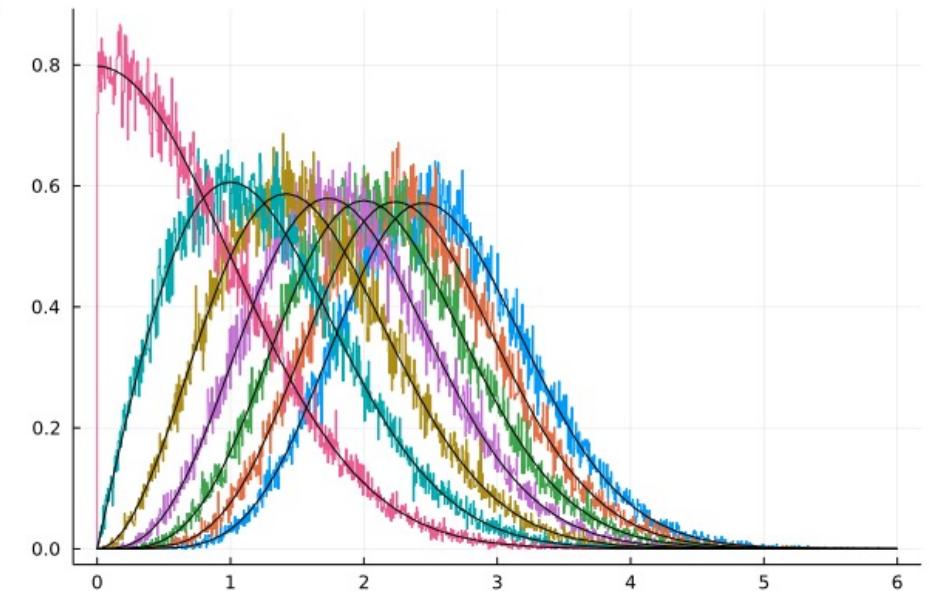
```
function qrhist():
    nsamples = 2^16,
    nbins = 1_000,
    xmin::Real = 0.0,
    xmax::Real = 6.0,
    matsize::Val{m} = Val(7),
    executor = nothing,
) where {m}
    hist = fill!(arrayon(executor, Int, (nbins, m)), 0)
    dx = (xmax - xmin) / nbins
    @floop executor for _i in 1:nsamples
        M = randn(SMatrix{m,m,Float32})
        _q, r = qr(M)
        for i in 1:m
            bin = floor(Int, (abs(r[i, i]) - xmin) / dx) + 1
            @inbounds @atomic :monotonic hist[max(begin, min(bin, end)), i] += 1
        end
    end
    edges = collect(xmin:dx:xmax)
    return [Histogram(edges, collect(h), :left, false) for h in eachcol(hist)]
end
```

Execution on multiple CPUs

```
In [4]: hist = normalize.(qrhist(nsamples = 2^16))

plt = plot(hist, seriestype = :step, legend = nothing)
plot_chi!(plt, length(hist))
```

Out[4]:



Write Once, Execute Anywhere!

```
function qrhist():
    nsamples = 2^18,
    nbins = 1_000,
    xmin::Real = 0.0,
    xmax::Real = 6.0,
    matsize::Val{m} = Val(7),
    executor = nothing,
) where {m}
    hist = fill!(arrayon(executor, Int, (nbins, m)), 0)
    dx = (xmax - xmin) / nbins
    @floop executor for _i in 1:nsamples
        M = randn(SMatrix{m,m,Float32})
        _q, r = qr(M)
        for i in 1:m
            bin = floor(Int, (abs(r[i, i]) - xmin) / dx) + 1
            @inbounds @atomic :monotonic hist[max(begin, min(bin, end)), i] += 1
        end
    end
    edges = collect(xmin:dx:xmax)
    return [Histogram(edges, collect(h), :left, false) for h in eachcol(hist)]
end
```

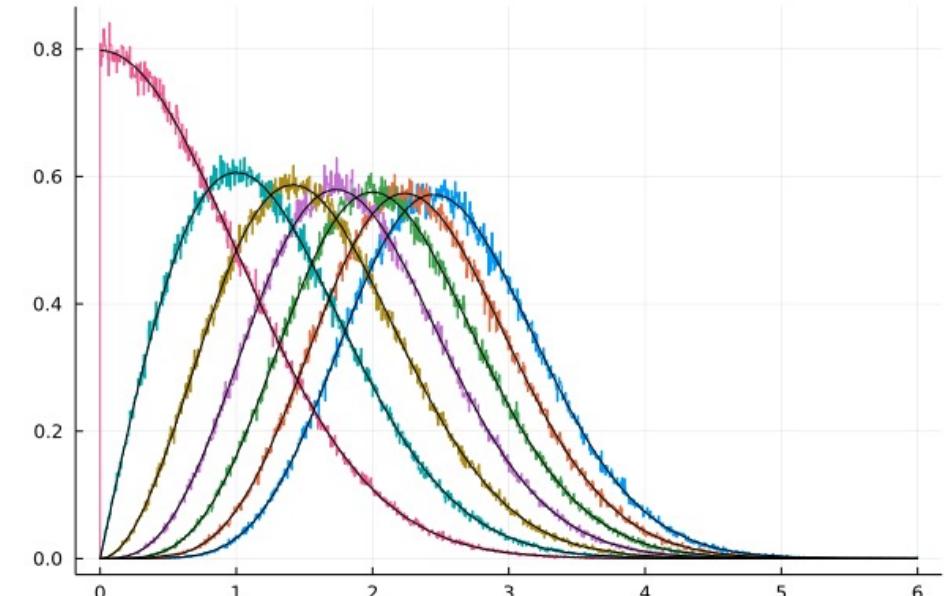
Execution on GPU

```
In [5]: using AtomicArraysCUDA
using CUDA
using FoldscUDA
arrayon(::CUDAEx, T, size) = CuArray{T}(undef, size)

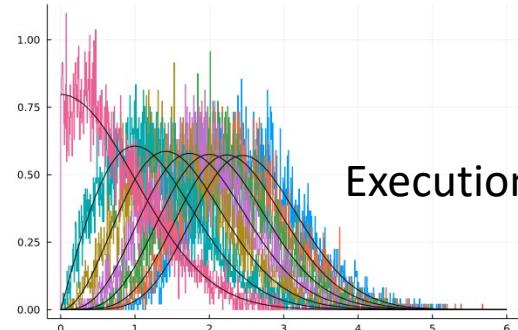
hist = normalize.(qrhist.executor = CUDAEx(), nsamples = 2^19))

plt = plot(hist, seriestype = :step)
plot_chi!(plt, length(hist))
```

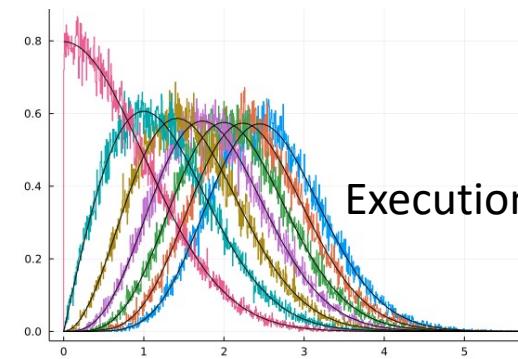
Out[5]:



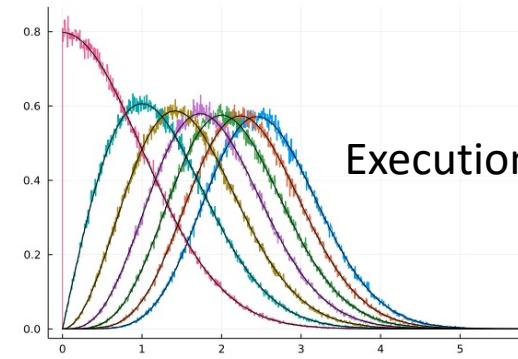
No sacrifice in Performance!



Execution on a single CPU



Execution on multiple CPUs



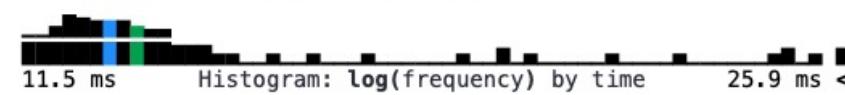
Execution on GPU

```
BenchmarkTools.Trial: 550 samples with 1 evaluation.  
Range (min ... max): 9.045 ms ... 10.760 ms | GC (min ... max): 0.00% ... 0.00%  
Time (median): 9.082 ms | GC (median): 0.00%  
Time (mean ± σ): 9.090 ms ± 76.303 μs | GC (mean ± σ): 0.00% ± 0.00%
```



Memory estimate: 118.73 KiB, allocs estimate: 18.

```
BenchmarkTools.Trial: 373 samples with 1 evaluation.  
Range (min ... max): 11.491 ms ... 59.282 ms | GC (min ... max): 0.00% ... 0.00%  
Time (median): 12.846 ms | GC (median): 0.00%  
Time (mean ± σ): 13.417 ms ± 3.219 ms | GC (mean ± σ): 0.00% ± 0.00%
```



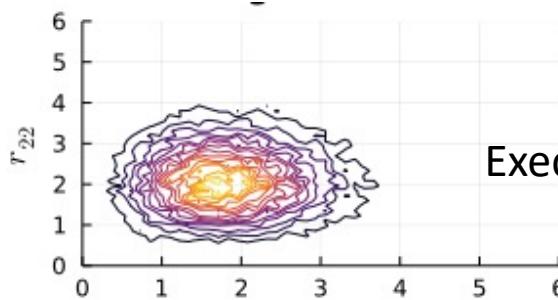
Memory estimate: 127.33 KiB, allocs estimate: 128.

```
BenchmarkTools.Trial: 517 samples with 1 evaluation.  
Range (min ... max): 9.202 ms ... 32.841 ms | GC (min ... max): 0.00% ... 0.00%  
Time (median): 9.487 ms | GC (median): 0.00%  
Time (mean ± σ): 9.655 ms ± 1.574 ms | GC (mean ± σ): 0.00% ± 0.00%
```

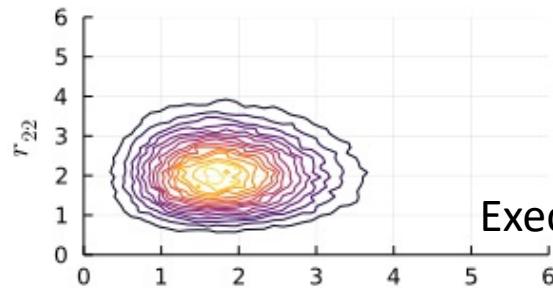


Memory estimate: 138.11 KiB, allocs estimate: 391.

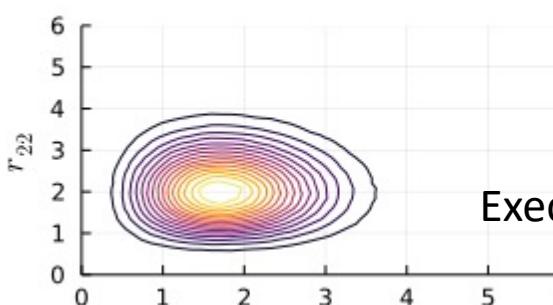
Even more prominent in 2D!



Execution on a single CPU

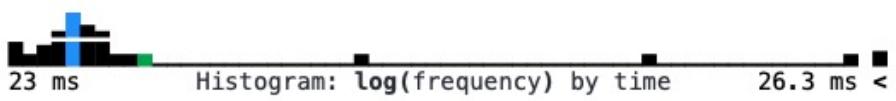


Execution on multiple CPUs



Execution on GPU

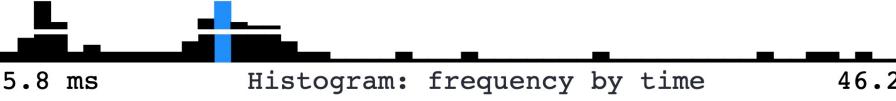
BenchmarkTools.Trial: 213 samples with 1 evaluation.
Range (min ... max): 22.973 ms ... 65.061 ms | GC (min ... max): 0.00% ... 64.04%
Time (median): 23.228 ms | GC (median): 0.00%
Time (mean ± σ): 23.474 ms ± 2.889 ms | GC (mean ± σ): 0.83% ± 4.39%



Histogram: log(frequency) by time
23 ms 26.3 ms <

Memory estimate: 488.42 KiB, allocs estimate: 3.

BenchmarkTools.Trial: 164 samples with 1 evaluation.
Range (min ... max): 25.780 ms ... 75.986 ms | GC (min ... max): 0.00% ... 0.00%
Time (median): 30.574 ms | GC (median): 0.00%
Time (mean ± σ): 30.560 ms ± 5.035 ms | GC (mean ± σ): 0.00% ± 0.00%



Histogram: frequency by time
25.8 ms 46.2 ms <

Memory estimate: 497.02 KiB, allocs estimate: 113.

BenchmarkTools.Trial: 185 samples with 1 evaluation.
Range (min ... max): 26.943 ms ... 27.240 ms | GC (min ... max): 0.00% ... 0.00%
Time (median): 27.067 ms | GC (median): 0.00%
Time (mean ± σ): 27.065 ms ± 43.939 µs | GC (mean ± σ): 0.00% ± 0.00%



Histogram: frequency by time
26.9 ms 27.2 ms <

Memory estimate: 19.84 KiB, allocs estimate: 354.

Productivity: Focus not on speedup but what can be accomplished / unit time

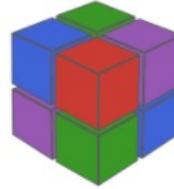
Showing Impact of Language

- Case Study: A Random matrix histogram
 - No extra hoops to jump through for parallelism and GPU acceleration (no CUDA, no MPI,...)
 - distribution of diagonals of $\text{qr}(\text{randn}(n,n))$

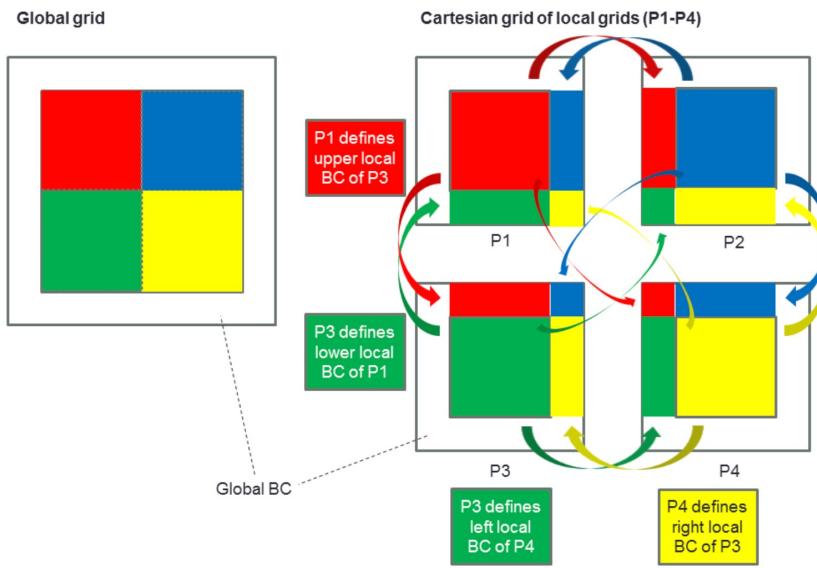
Showing Impact of Language

- Case Study: A Random matrix histogram
 - No extra hoops to jump through for parallelism and GPU acceleration (no CUDA, no MPI,...)
 - distribution of diagonals of $\text{qr}(\text{randn}(n,n))$
- Demo: Diffusion
 - No Reinventing the Wheel – Build upon solid foundation
 - You have a C/MPI or Fortran/MPI Code? No need to rewrite
 - You need a ghost/halo code? No need to write
 - Can easily compose parallel modes
 - Message Passing and Task Parallelism

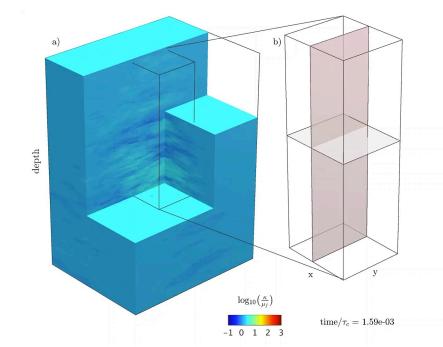
ImplicitGlobalGrid



- Samuel Omlin & Ludovic Räss Swiss Computing Center
- The most common HPC request: stencils & halo/ghost cells
- Made nearly trivial!



5



Spontaneous formation of fluid escape pipes from subsurface reservoirs

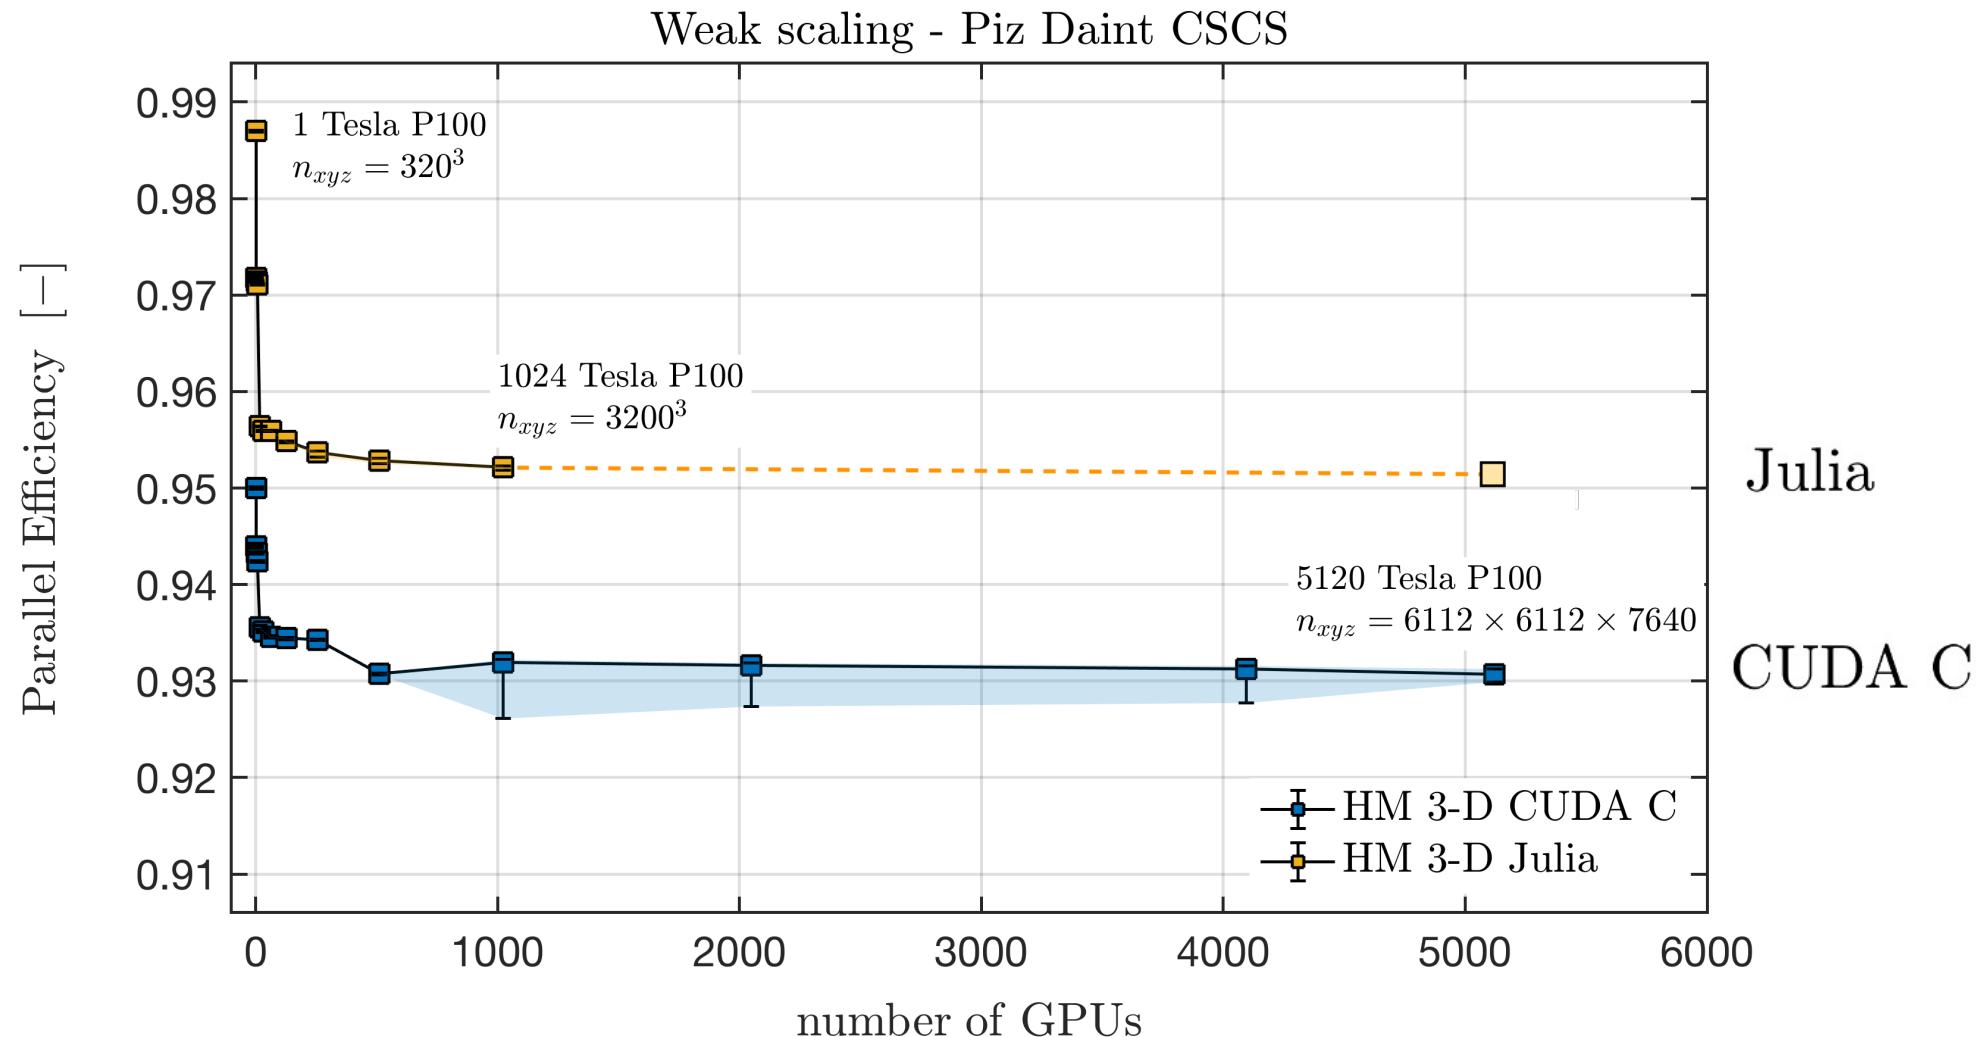
Ludovic Räss Nina S. C. Simon & Yury Y. Podladchikov

Hide communication:

- Compute boundary region and interior points on two different streams.
- As soon as boundary region has finished, start halo update

Reported Results

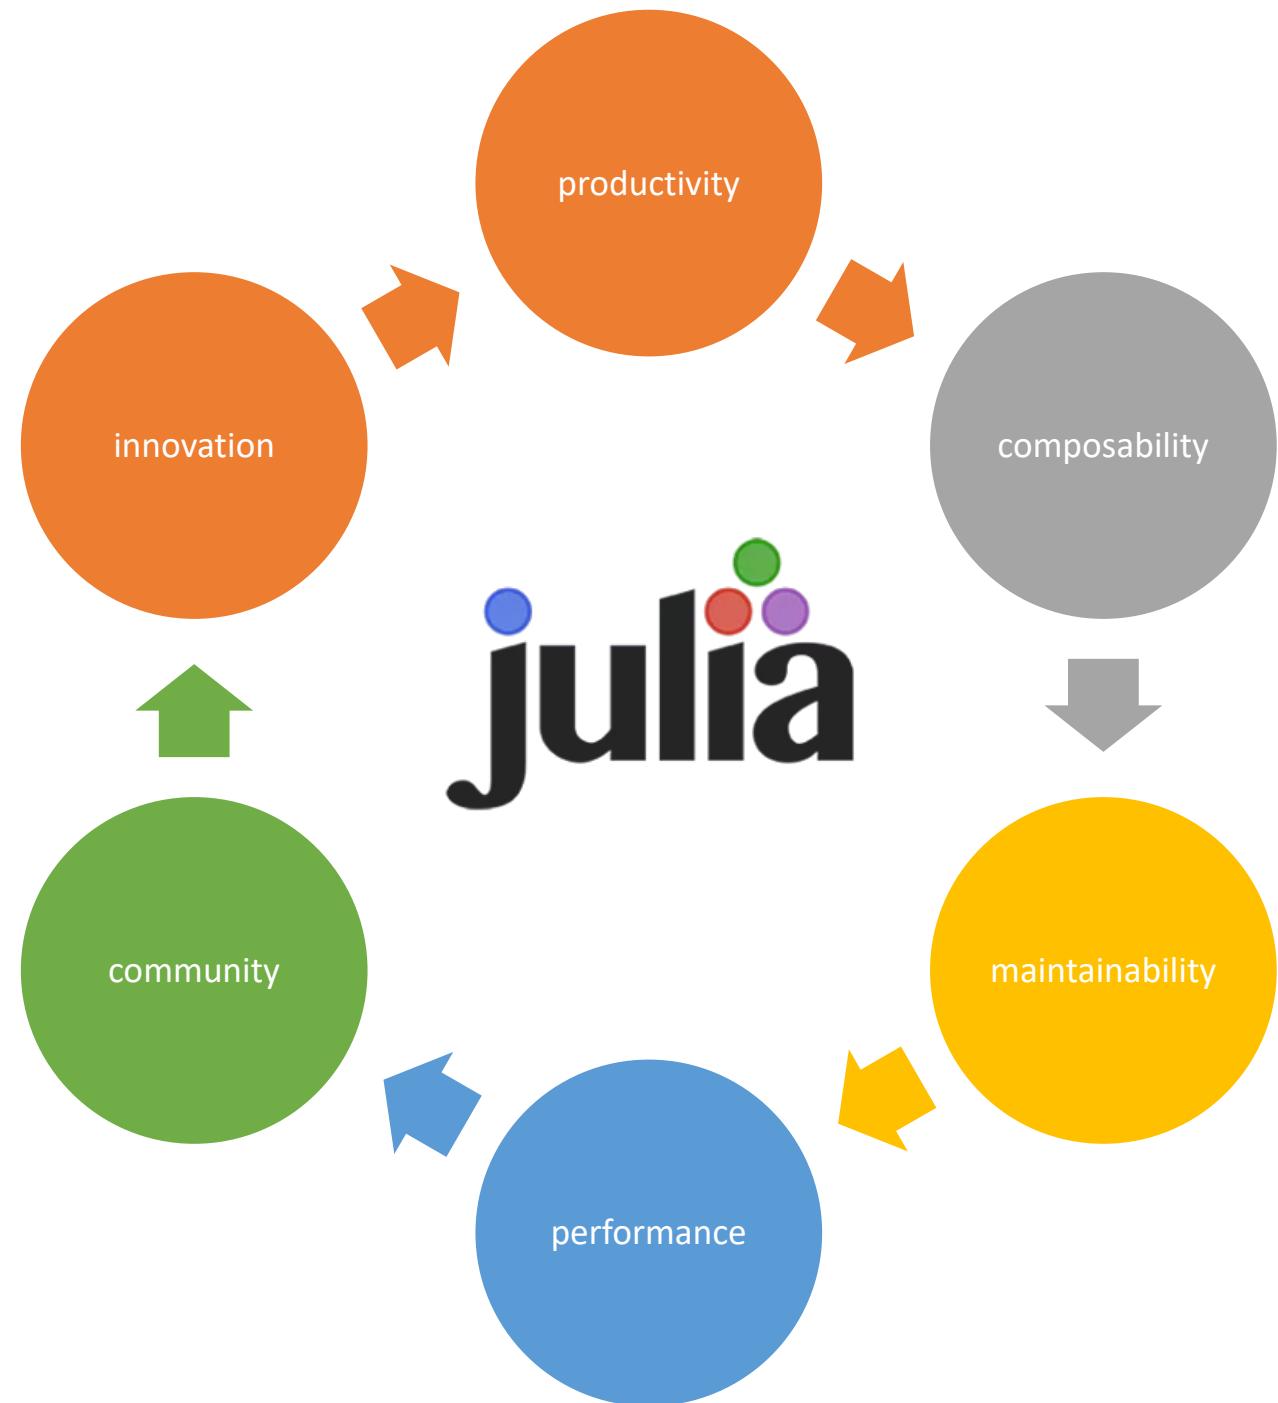
(ImplicitGlobalGrid.jl github)



Demos: Take home message

- You have a C/MPI or Fortran/MPI? No need to rewrite!
- You need a ghost / halo code? No need to rewrite!
- You want to compose different parallel modes? Easy!

No re-inventing the wheel – Build upon solid foundations



Solving Big Challenges in our times

Building and deploying a scalable, composable workflow for CESMIX in Julia

How do provide
binaries for **all**
platforms?

How do we make
reproducibility
simple?

How do we
empower user?

How do we support
cross-language AD?

How do we allocate
resources and
schedule
subsystems?

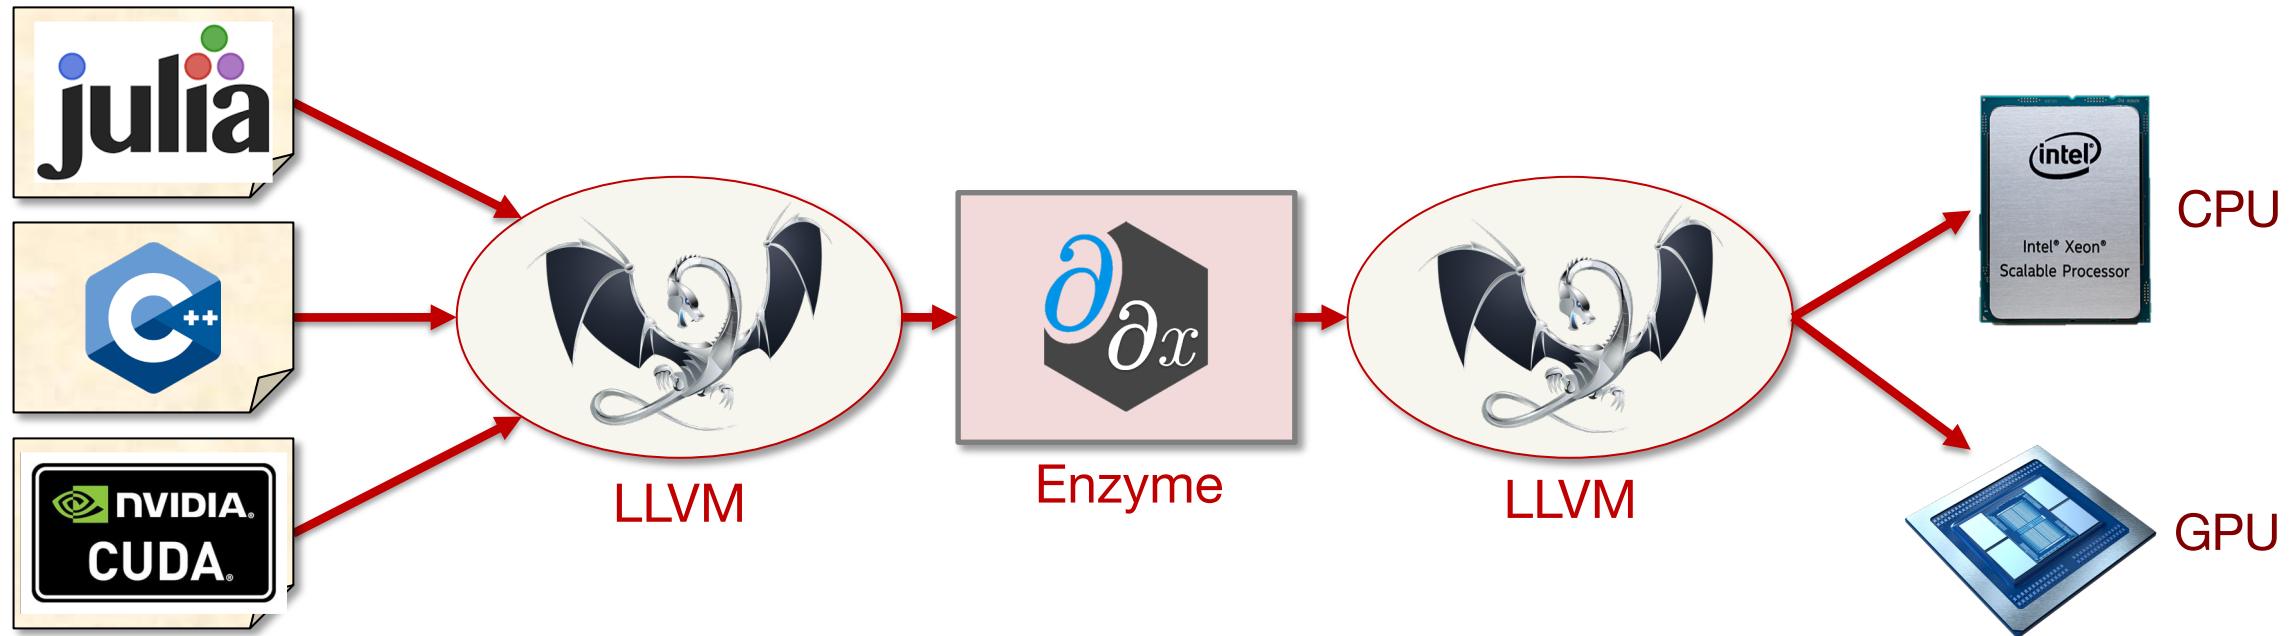
How do we **not**
accidentally DDOS a
cluster?

Valentin Churavy
JuliaLab, CSAIL
October 2021

vchuravy@mit.edu

Computing derivatives: Enzyme [MC20]

Enzyme performs automatic differentiation (AD) **within** the compiler.



Enzyme differentiates through
MPI.

MIT News: New climate model to be built from the ground up

MIT News
ON CAMPUS AND AROUND THE WORLD

December 12, 2018

To take advantage of new computer architectures, languages, and machine learning techniques, the team has partnered with [Alan Edelman](#)'s group in CSAIL at MIT, who will help write the new generation climate model in the Julia computing language developed by the group. This will enable the MIT team to target GPUs, CPUs, and evolving computer architectures within one code base.



Tapio Schneider



Chris Hill

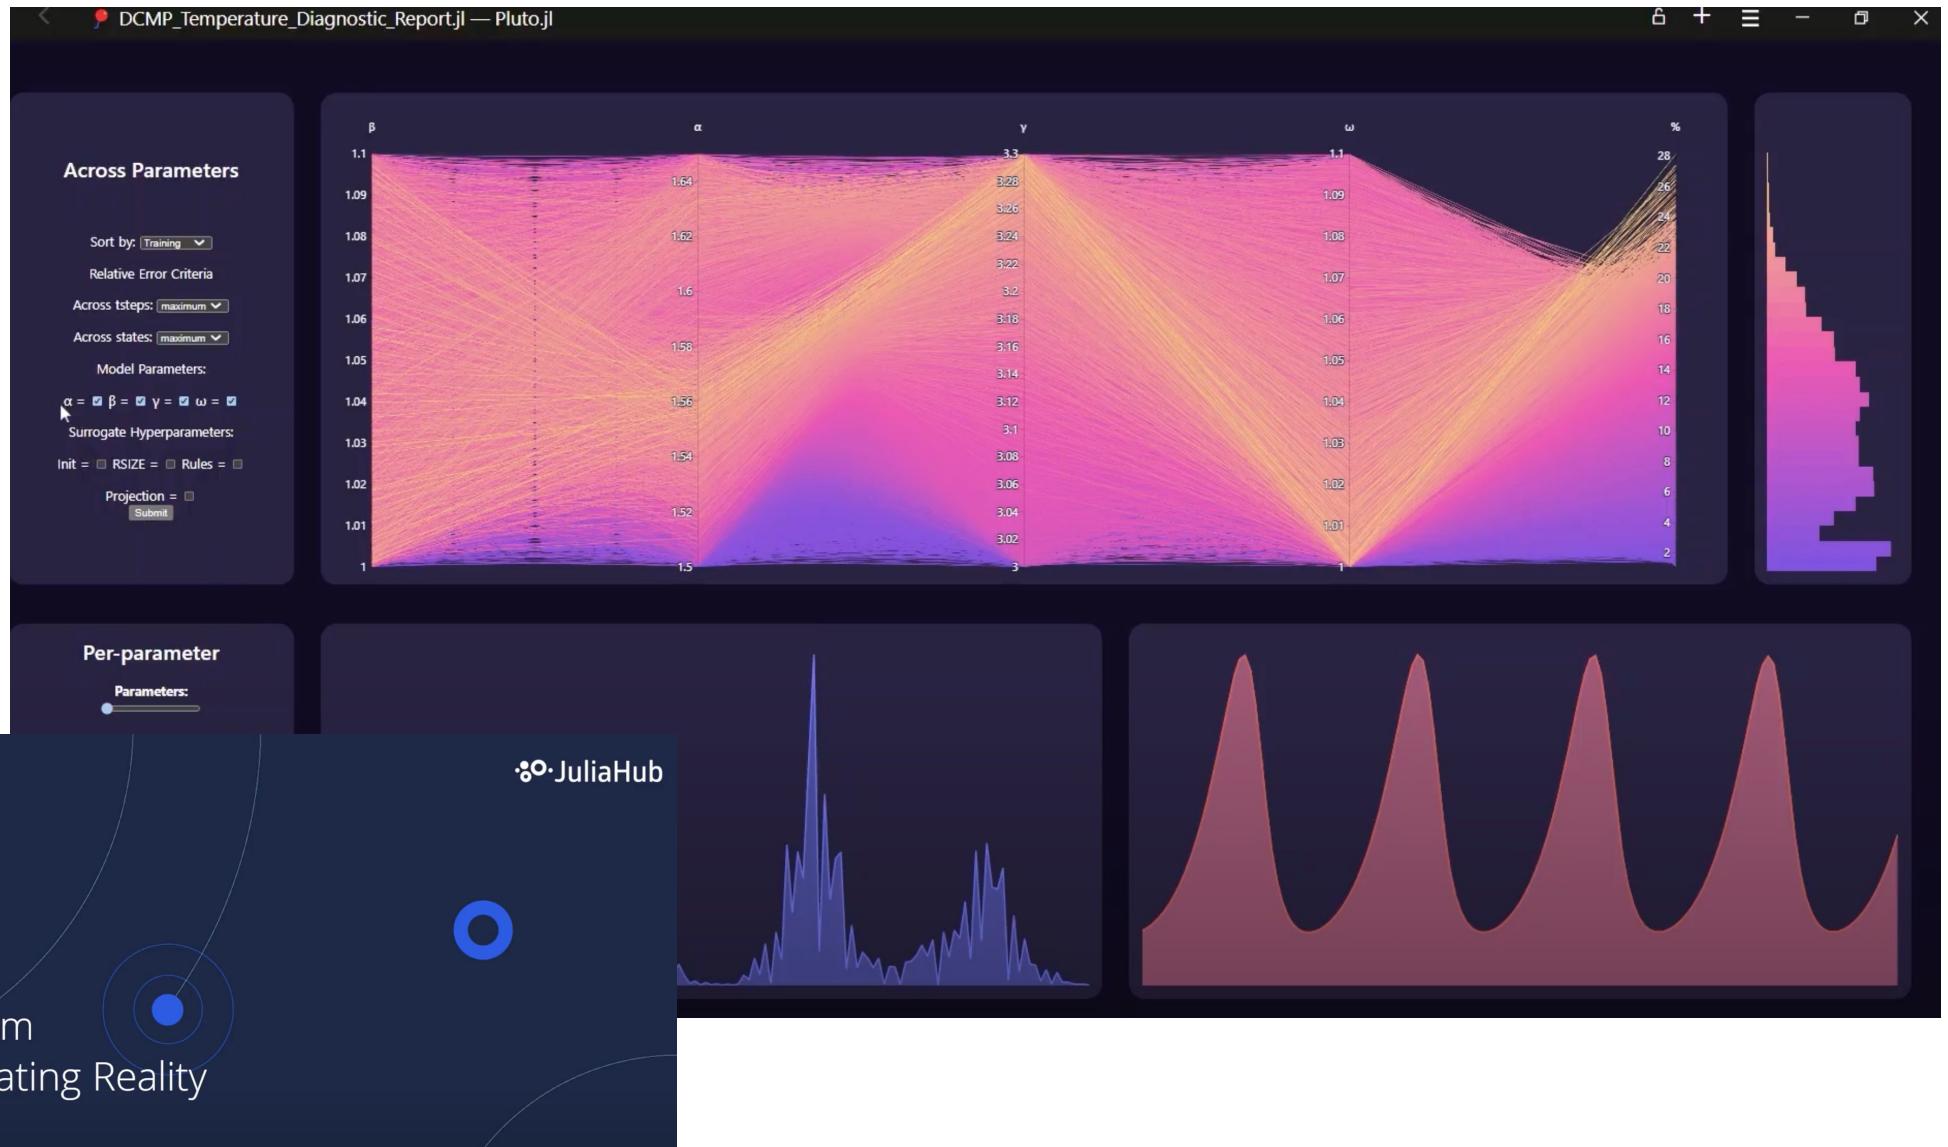


Raffaele Ferrari



John Marshall

JuliaSim Simulating Reality



Code is the “New” Math

```
function volumerhs(Q:NamedTuple{S, NTuple{3, T}}; bathymetry, metric, b, u, elem, gravity, hnl) where {S, T}
    (rhs_h, rhs_u, rhs_V) = (rhs_h, rhs_u, rhs_V)
    (S, u, V) = (Q.S, Q.u, Q.V)
    Ng = size(b, 1)
    metric.J
    dim = 3
    (x, y) = metric(x, metric.y)
    (x, y, z) = metric(x, metric.y, metric.z)
    fluxArray(Oflast,:)\converg(dim,Ng,Ng)
    fluxArray(Oflast,:)\converg(dim,Ng,Ng)
    fluxArray(Oflast,:)\converg(dim,Ng,Ng)
    for e in 1:Ng
        #Get primitive variables and fluxes
        h= bathymetry[1,e]
        hb = bathymetry[2,e]
        hnc = h + hb
        uQ[1,:,e] = J_h
        vQ[1,:,e] = J_v
        fluxh[1,:,e]\Q[1,:,e]
        fluxh[2,:,e]\Q[1,:,e] = 0.5 * gravity * ha.^2 * sal * gravity * ha.^2 * hb
        fluxh[2,:,e]\Q[1,:,e] = 0.5 * gravity * ha.^2 * sal
        fluxv[1,:,e]\Q[1,:,e] = 0.5 * gravity * ha.^2 * sal
        fluxv[1,:,e]\Q[1,:,e] = 0.5 * gravity * ha.^2 * sal
        fluxv[2,:,e]\Q[1,:,e] = 0.5 * gravity * ha.^2 * sal * gravity * ha.^2 * hb
        #Loop over n-grid lines
        for j = 1:Ng
            rhah[1,:,e] \= D + (d11 * u.^2)^(1,0,0) * (Nx[1,:,e]) * fluxh[1,:,e] + (d11 * v.^2)^(1,0,0) * fluxh[1,:,e]
            rhah[1,:,e] \= D + (d11 * u.^2)^(1,0,0) * (Nx[1,:,e]) * fluxh[1,:,e] + (d11 * v.^2)^(1,0,0) * fluxh[1,:,e]
            rhav[1,:,e] \= D + (d11 * u.^2)^(1,0,0) * (Ny[1,:,e]) * fluxv[1,:,e] + (d11 * v.^2)^(1,0,0) * fluxv[1,:,e]
            rhav[1,:,e] \= D + (d11 * u.^2)^(1,0,0) * (Ny[1,:,e]) * fluxv[1,:,e] + (d11 * v.^2)^(1,0,0) * fluxv[1,:,e])
        end #for j
    end #for e
end #function volumerhs_2d
```

$$\begin{aligned}\Delta h[\text{elem}] &+= \int \nabla \Psi(d\vec{x} * \vec{U}_e * \mathbf{J}) \\ \Delta \vec{U}[\text{elem}] &+= \int \nabla \Psi(d\vec{x} * (\vec{U}_e * \vec{U}_e' / ht_e + g * (ht_e^2 - hb_e^2)/2 * \mathbf{I}) * \mathbf{J})\end{aligned}$$

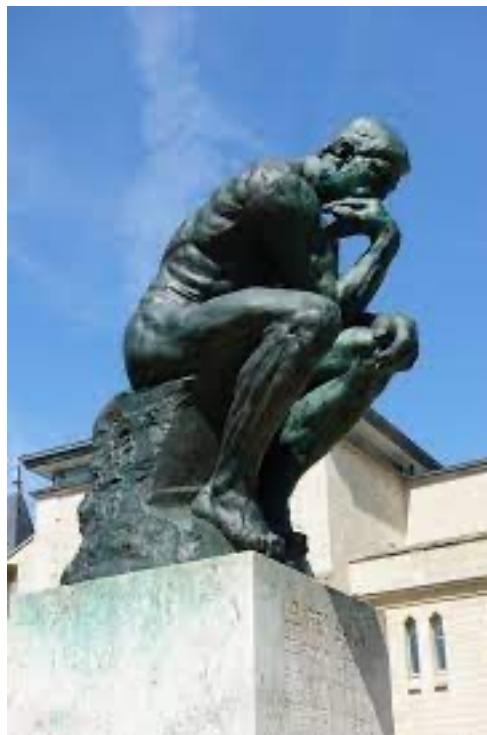
$$\int_{\Omega_e} \nabla \psi \cdot \mathbf{f}_N^{(e)} d\Omega_e$$

$$\mathbf{f} = \left(\mathbf{U}, \frac{\mathbf{U} \otimes \mathbf{U}}{h} + g(h^2 - h_b^2) \mathbf{I}_2 \right)$$

You don’t “code” the “math,” “code” is the “math”

Conclusion

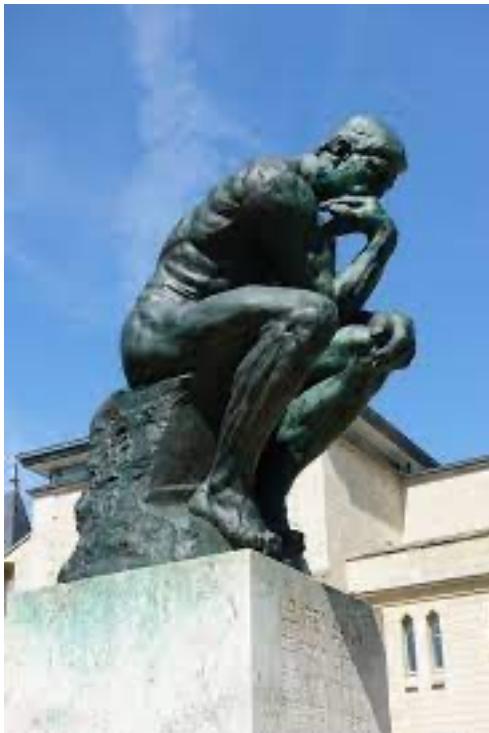
SIAM Conference on Parallel Processing for Scientific Computing (PP22)



Making Parallel Computing Easier

Conclusion

SIAM Conference on Parallel Processing for Scientific Computing (PP22)



Making Parallel Computing Easier
Join the great community:

julia